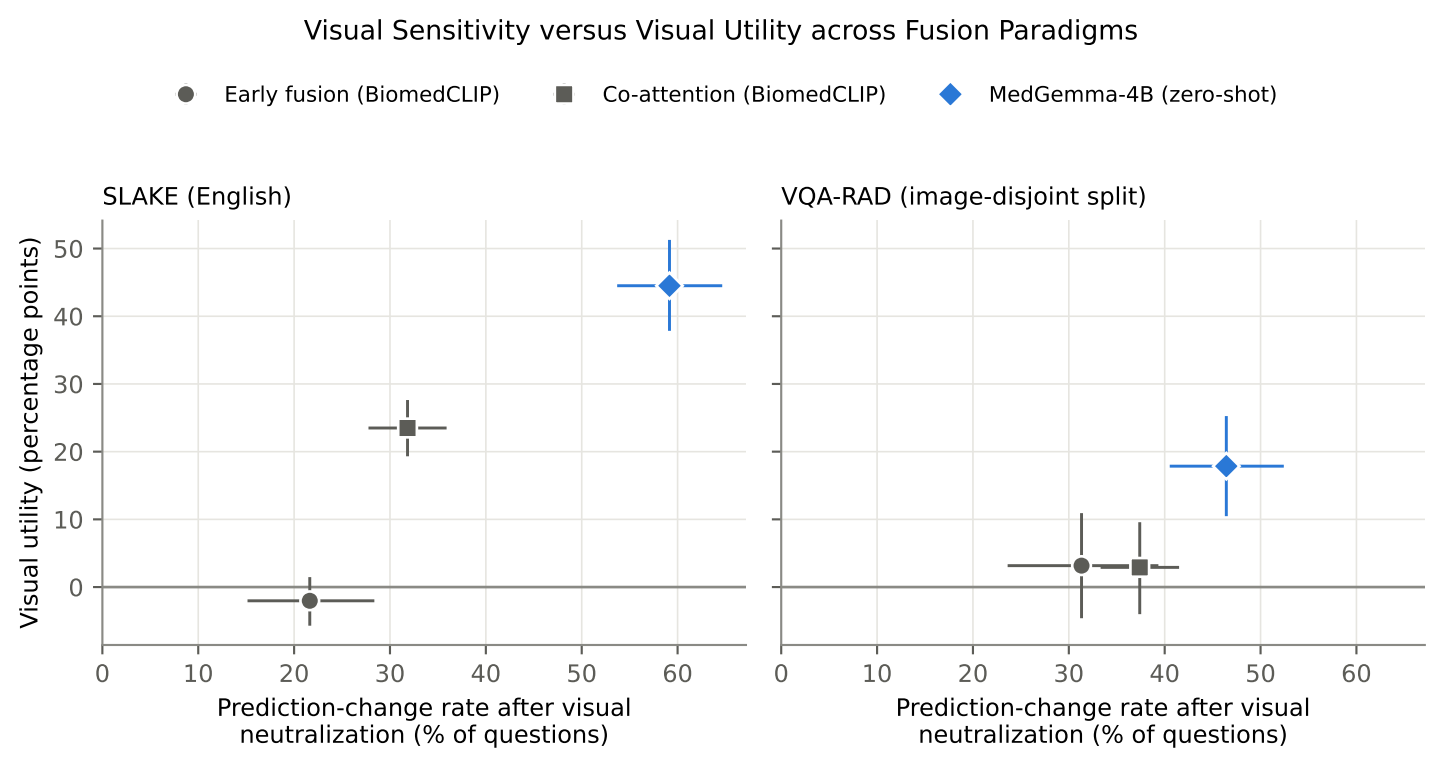
Visual Sensitivity versus Visual Utility across Fusion Paradigms
Early fusion (BiomedCLIP)
Co-attention (BiomedCLIP)
MedGemma-4B (zero-shot)
SLAKE (English)
VQA-RAD (image-disjoint split)
Visual utility (percentage points)
50
40
30
20
10
0
0
10
20
30
40
50
60
Prediction-change rate after visual neutralization (% of questions)
Prediction-change rate after visual neutralization (% of questions)